

Name : Tushar Shirsath

Roll No : 220940325083

Q.1

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv(r'C:\Users\Dell\Desktop\ML_Module_Exam\Data\car.csv')
df
```

```
Out[2]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0
...
296	city	2016	9.50	11.60	33988	Diesel	Dealer	Manual	0
297	brio	2015	4.00	5.90	60000	Petrol	Dealer	Manual	0
298	city	2009	3.35	11.00	87934	Petrol	Dealer	Manual	0
299	city	2017	11.50	12.50	9000	Diesel	Dealer	Manual	0
300	brio	2016	5.30	5.90	5464	Petrol	Dealer	Manual	0

301 rows × 9 columns

1.Data understanding and exploration

```
In [3]: df.shape
```

```
Out[3]: (301, 9)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Kms_Driven      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Seller_Type     301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

```
In [6]: df.describe(include='O')
```

```
Out[6]:
```

	Car_Name	Fuel_Type	Seller_Type	Transmission
count	301	301	301	301
unique	98	3	2	2
top	city	Petrol	Dealer	Manual
freq	26	239	195	261

```
In [7]: df.isnull().sum()
```

```
Out[7]: Car_Name      0
Year      0
Selling_Price  0
Present_Price  0
Kms_Driven   0
Fuel_Type    0
Seller_Type   0
Transmission  0
Owner        0
dtype: int64
```

```
In [8]: df.nunique()
```

```
Out[8]: Car_Name      98
Year      16
Selling_Price  156
Present_Price  147
Kms_Driven   206
Fuel_Type     3
Seller_Type   2
Transmission  2
Owner        3
dtype: int64
```

```
In [10]: df.shape
```

```
Out[10]: (301, 9)
```

```
In [11]: df.isnull().sum()/len(df) * 100
```

```
Out[11]: Car_Name      0.0
Year      0.0
Selling_Price  0.0
Present_Price  0.0
Kms_Driven   0.0
Fuel_Type    0.0
Seller_Type   0.0
Transmission  0.0
Owner        0.0
dtype: float64
```

```
In [ ]: # No null values are present
```

```
In [13]: df.columns
```

```
Out[13]: Index(['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Kms_Driven',
               'Fuel_Type', 'Seller_Type', 'Transmission', 'Owner'],
              dtype='object')
```

```
In [12]: # Car Name can be drop
```

```
In [14]: df.drop(['Car_Name'], axis=1, inplace=True)
```

```
In [15]: df.shape
```

```
Out[15]: (301, 8)
```

```
In [16]: df.Year.value_counts()
```

```
Out[16]: 2015    61
         2016    50
         2014    38
         2017    35
         2013    33
         2012    23
         2011    19
         2010    15
         2008     7
         2009     6
         2006     4
         2005     4
         2003     2
         2007     2
         2018     1
         2004     1
         Name: Year, dtype: int64
```

```
In [17]: df.Year.nunique()
```


```
Out[17]: 16
```

```
In [18]: sns.boxplot(df.Year)
```

C:\Users\Dell\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[18]: <AxesSubplot:xlabel='Year'>
```



```
In [19]: df.columns
```

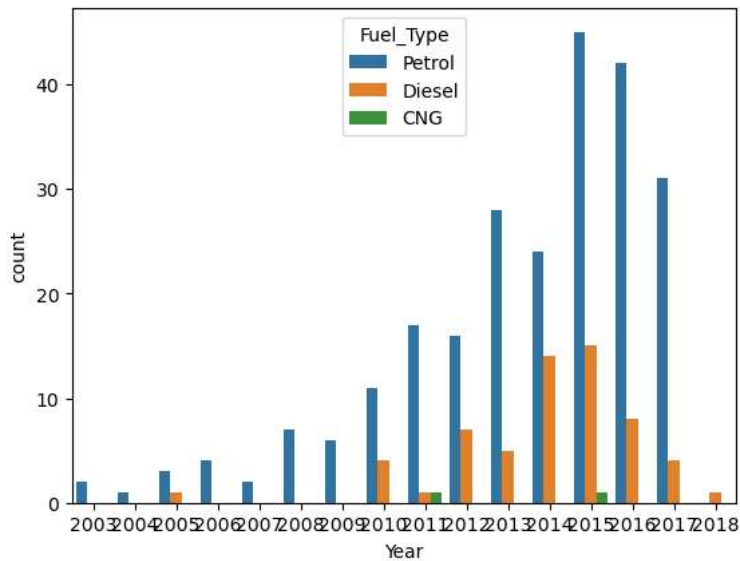
```
Out[19]: Index(['Year', 'Selling_Price', 'Present_Price', 'Kms_Driven', 'Fuel_Type',
              'Seller_Type', 'Transmission', 'Owner'],
              dtype='object')
```

```
In [20]: # Selling price vs Year
```

```
In [25]: sns.countplot('Year', hue = 'Fuel_Type', data=df)
```

C:\Users\Dell\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword argument: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()


```
Out[25]: <AxesSubplot:xlabel='Year', ylabel='count'>
```



```
In [21]: sns.regplot('Year', 'Selling_Price', data = df)
```

C:\Users\Dell\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword arguments: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()

```
Out[21]: <AxesSubplot:xlabel='Year', ylabel='Selling_Price'>
```



```
In [ ]: # As the year increases selling price also increases
```


```
In [22]: df.Selling_Price.describe()
```

```
Out[22]: count    301.000000
mean        4.661296
std         5.082812
min         0.100000
25%         0.900000
50%         3.600000
75%         6.000000
max         35.000000
Name: Selling_Price, dtype: float64
```

```
In [23]: # Selling Price vs Present Price
```

```
In [24]: df.plot.scatter(x='Selling_Price', y='Present_Price')
```

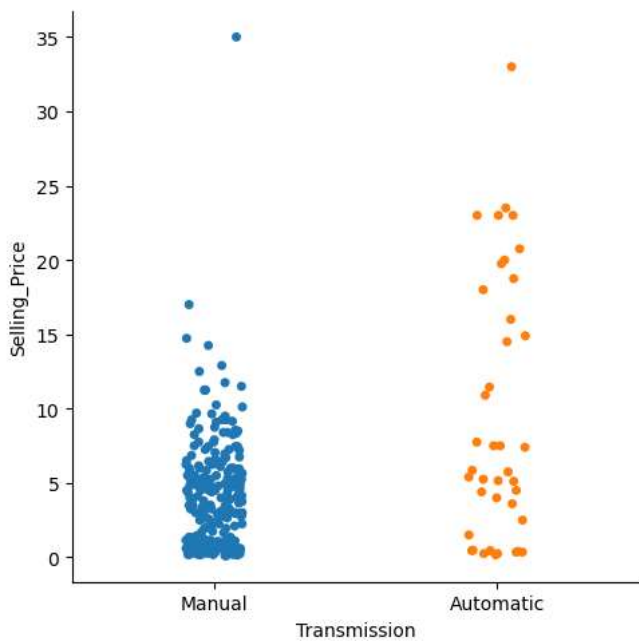
```
Out[24]: <AxesSubplot: xlabel='Selling_Price', ylabel='Present_Price'>
```



```
In [ ]:
```

```
In [26]: sns.catplot(data=df, x='Transmission', y='Selling_Price')
```


```
Out[26]: <seaborn.axisgrid.FacetGrid at 0x195c7447d30>
```



```
In [ ]: # Automatic transmission has high selling price as compared to Manual Transmission
```

```
In [27]: df.plot.scatter('Selling_Price', 'Kms_Driven')
```


```
Out[27]: <AxesSubplot:xlabel='Selling_Price', ylabel='Kms_Driven'>
```



```
In [ ]:
```

```
In [30]: sns.boxplot(x='Year', y='Selling_Price', data=df)
```

```
Out[30]: <AxesSubplot:xlabel='Year', ylabel='Selling_Price'>
```



2.Data cleaning

```
In [31]: df.columns
```

```
Out[31]: Index(['Year', 'Selling_Price', 'Present_Price', 'Kms_Driven', 'Fuel_Type',  
              'Seller_Type', 'Transmission', 'Owner'],  
              dtype='object')
```

```
In [32]: df.head()
```

```
Out[32]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

In [33]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Year                  301 non-null    int64
1   Selling_Price         301 non-null    float64
2   Present_Price         301 non-null    float64
3   Kms_Driven            301 non-null    int64
4   Fuel_Type             301 non-null    object
5   Seller_Type           301 non-null    object
6   Transmission          301 non-null    object
7   Owner                 301 non-null    int64
dtypes: float64(2), int64(3), object(3)
memory usage: 18.9+ KB
```

In [34]: df.Year.value_counts()

```
Out[34]: 2015    61
         2016    50
         2014    38
         2017    35
         2013    33
         2012    23
         2011    19
         2010    15
         2008     7
         2009     6
         2006     4
         2005     4
         2003     2
         2007     2
         2018     1
         2004     1
Name: Year, dtype: int64
```

In [35]: *# Converting year column into number of years selling car is old*

```
In [36]: df['New_Year'] = 2022
df['Years'] = df.New_Year - df.Year
```

In [37]: df.head(1)

```
Out[37]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	New_Year	Years
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	2022	8

In [38]: df.drop(['Year', 'New_Year'], axis=1, inplace = True)

In [39]: df.head(1)

```
Out[39]:
```

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Years
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	8

In [40]: df.shape

Out[40]: (301, 8)

```
In [41]: df = pd.get_dummies(df, drop_first=True)
df.head()
```

```
Out[41]:
```

	Selling_Price	Present_Price	Kms_Driven	Owner	Years	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission_Manual
0	3.35	5.59	27000	0	8	0	1	0	1
1	4.75	9.54	43000	0	9	1	0	0	1
2	7.25	9.85	6900	0	5	0	1	0	1
3	2.85	4.15	5200	0	11	0	1	0	1
4	4.60	6.87	42450	0	8	1	0	0	1

In [42]: df.shape

Out[42]: (301, 9)

In [43]: *# sns.pairplot(data=df)*


In [44]: `df.corr()`

Out[44]:

	Selling_Price	Present_Price	Kms_Driven	Owner	Years	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission_
Selling_Price	1.000000	0.878983	0.029187	-0.088344	-0.236141	0.552339	-0.540571	-0.550724	-0.
Present_Price	0.878983	1.000000	0.203647	0.008057	0.047584	0.473306	-0.465244	-0.512030	-0.
Kms_Driven	0.029187	0.203647	1.000000	0.089216	0.524342	0.172515	-0.172874	-0.101419	-0.
Owner	-0.088344	0.008057	0.089216	1.000000	0.182104	-0.053469	0.055687	0.124269	-0.
Years	-0.236141	0.047584	0.524342	0.182104	1.000000	-0.064315	0.059959	0.039896	-0.
Fuel_Type_Diesel	0.552339	0.473306	0.172515	-0.053469	-0.064315	1.000000	-0.979648	-0.350467	-0.
Fuel_Type_Petrol	-0.540571	-0.465244	-0.172874	0.055687	0.059959	-0.979648	1.000000	0.358321	0.
Seller_Type_Individual	-0.550724	-0.512030	-0.101419	0.124269	0.039896	-0.350467	0.358321	1.000000	0.
Transmission_Manual	-0.367128	-0.348715	-0.162510	-0.050316	-0.000394	-0.098643	0.091013	0.063240	1.

In [45]: `sns.heatmap(df.corr(), annot=True)`

Out[45]: <AxesSubplot:>



3.Data preparation for Model Building

In [46]: `x = df.drop('Selling_Price', axis=1)`
`y = df['Selling_Price']`

In [47]: x

```
Out[47]:
```

	Present_Price	Kms_Driven	Owner	Years	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission_Manual
0	5.59	27000	0	8	0	1	0	1
1	9.54	43000	0	9	1	0	0	1
2	9.85	6900	0	5	0	1	0	1
3	4.15	5200	0	11	0	1	0	1
4	6.87	42450	0	8	1	0	0	1
...
296	11.60	33988	0	6	1	0	0	1
297	5.90	60000	0	7	0	1	0	1
298	11.00	87934	0	13	0	1	0	1
299	12.50	9000	0	5	1	0	0	1
300	5.90	5464	0	6	0	1	0	1

301 rows × 8 columns

In [48]: y

```
Out[48]:
```

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	...
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling_Price, Length: 301, dtype: float64

In [49]: x.shape, y.shape

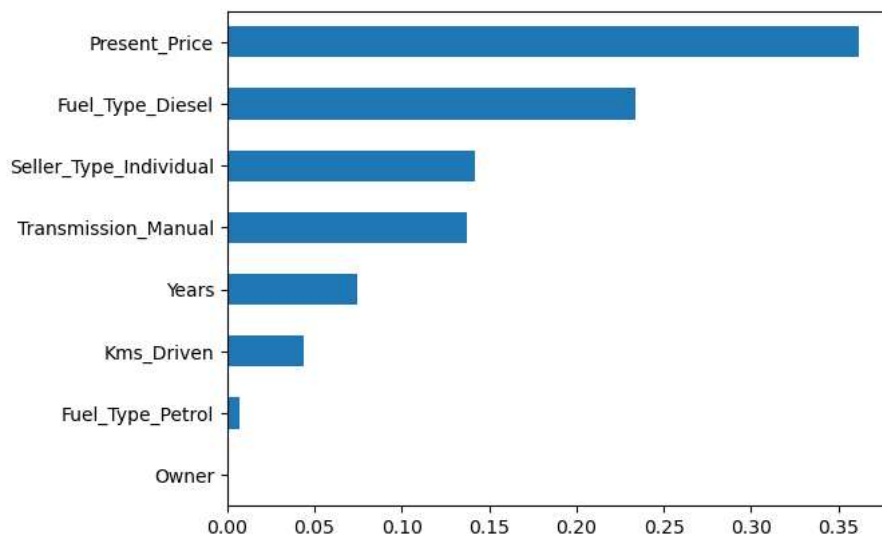
Out[49]: ((301, 8), (301,))

In []:

```
In [50]: from sklearn.ensemble import ExtraTreesRegressor
etr = ExtraTreesRegressor()
etr.fit(x,y)
```

Out[50]: ExtraTreesRegressor()

```
In [51]: feature = pd.Series(etr.feature_importances_, index=x.columns).sort_values(ascending=True)
feature.plot(kind = 'barh')
plt.show()
```



In []:

```
In [52]: from sklearn.model_selection import train_test_split
```

```
In [53]: x_train, x_test, y_train, y_test = train_test_split(x,y,train_size=0.8, random_state=42)
```

```
In [54]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
Out[54]: ((240, 8), (61, 8), (240,), (61,))
```

```
In [ ]:
```

```
In [55]: from sklearn.preprocessing import MinMaxScaler
```

```
In [56]: sc = MinMaxScaler()
```

```
In [57]: x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
In [58]: x_train, x_test
```

```
Out[58]: (array([[0.00465973, 0.05105105, 0.33333333, ..., 1.          , 1.          ,
        1.          ],
       [0.00682705, 0.00600601, 0.          , ..., 1.          , 1.          ,
        1.          ],
       [0.00506068, 0.0990991 , 0.          , ..., 1.          , 1.          ,
        1.          ],
       ...,
       [0.03391851, 0.03203203, 0.33333333, ..., 1.          , 1.          ,
        1.          ],
       [0.10489814, 0.13781982, 0.          , ..., 1.          , 0.          ,
        1.          ],
       [0.01582141, 0.00700701, 0.          , ..., 1.          , 1.          ,
        1.          ]]),
 array([[ 0.00270915,  0.04704705,  0.          ,  0.07142857,  0.          ,
        1.          ,  1.          ,  0.          ],
       [ 0.14390984,  0.02098098,  0.          ,  0.07142857,  0.          ,
        1.          ,  0.          ,  1.          ],
       [ 0.09839619,  0.11911912,  0.          ,  0.35714286,  1.          ,
        0.          ,  0.          ,  1.          ],
       ...,
       [ 0.00270915,  0.04704705,  0.          ,  0.07142857,  0.          ,
        1.          ,  1.          ,  0.          ]])
```

```
In [ ]:
```

4.Model building and evaluation

Random Forest Regressor

```
In [63]: from sklearn.ensemble import RandomForestRegressor
```

```
In [64]: model = RandomForestRegressor()
model.fit(x_train, y_train)
model.score(x_test, y_test)
```


```
Out[64]: 0.9590803492264692
```

```
In [71]: y_pred = model.predict(x_test)
y_pred
```

```
Out[71]: array([ 0.4437, 10.9463,  4.9125,  0.2212,  7.5805,  6.41   ,  1.0574,
        0.57   ,  0.4713,  6.8295,  7.9604,  1.0773,  8.1506,  0.4491,
        5.44   ,  2.6155,  1.158 , 14.4091,  0.4771,  1.5455,  0.3337,
        7.9959,  4.85   ,  2.746 ,  0.5072,  3.4855,  5.3825,  3.139 ,
        1.2138,  1.1585,  0.4162,  9.9475,  0.4603,  2.5245,  7.7289,
        4.2315,  6.1575,  6.036 ,  2.5585,  6.504 ,  4.1598,  3.4018,
        4.971 ,  0.5703,  6.2125,  0.7568,  8.2245,  7.3145,  2.9165,
        3.6135,  5.02   ,  1.552 , 21.6622, 19.8743,  6.472 , 10.2125,
        5.066 ,  8.8477,  2.593 ,  6.7579,  0.2407])
```

```
In [72]: plt.scatter(y_pred, y_test, color = 'Blue')
plt.plot(y_pred, y_test, color = 'red')
```

```
Out[72]: [matplotlib.lines.Line2D at 0x195cbcd0550>]
```



```
In [ ]:
```

5.Result with error calculation

```
In [73]: from sklearn import metrics
```

```
In [79]: # Mean Absolute error
round(metrics.mean_absolute_error(y_test, y_pred),2)
```

```
Out[79]: 0.64
```

```
In [80]: # Mean Squared error
round(metrics.mean_squared_error(y_test, y_pred),2)
```

```
Out[80]: 0.94
```

```
In [81]: # Median Absolute Error
round(metrics.median_absolute_error(y_test, y_pred),2)
```

```
Out[81]: 0.43
```

```
In [82]: # Explain Variance Factor
round(metrics.explained_variance_score(y_test, y_pred),2)
```

```
Out[82]: 0.96
```

```
In [83]: # Model Score
model.score(x_test, y_test)
```

```
Out[83]: 0.9590803492264692
```

```
In [84]: # Model r2- score
metrics.r2_score(y_test, y_pred)
```

```
Out[84]: 0.9590803492264692
```

```
In [ ]:
```

Conclusion

- For model building, top 3 important features are - 1.Present_Price, 2.Fuel Type Diesel, 3.Seller Type Individual
- I have got accuracy upto 95 % using Random Forest Regressor Model
- Automatic transmission has high selling price as compared to Manual Transmission
- As the year increases selling price also increases

```
In [ ]:
```

```
In [ ]:
```

